

Predicting Delivery Time with AI: A Full-Stack Approach to Revolutionize Hyperlocal Logistics

By [Your Name]

1. Introduction: Why Delivery Time Prediction is the Next Frontier

In today's hyper-competitive delivery ecosystem, speed isn't a luxury — it's an expectation. Whether it's piping hot pizza from Domino's or a food order from Zomato or Swiggy, consumers expect on-time deliveries. But ensuring this level of precision isn't easy. Urban congestion, unpredictable traffic, vehicle types, and human inefficiencies all affect the estimated time of arrival (ETA).

And while many companies rely on hard-coded heuristics or rigid delivery rules, there's a better way to do it: **predictive machine learning models** trained on delivery history and contextual signals.

This article presents a production-ready, full-stack solution to predicting delivery durations using AI. We'll walk through every step — from data generation to modeling, deployment, and business impact — to show how companies can leverage this system to enhance efficiency, customer satisfaction, and operational performance.

2. The Business Problem

Let's consider the pain points for a delivery company:

- Dispatching the wrong vehicle leads to longer travel times.
- Underestimating delivery time causes missed SLAs and penalties.
- Overestimating it leads to under-utilized resources.
- Lack of real-time, personalized ETA estimation.

A poor delivery estimation pipeline doesn't just impact operations; it affects customer trust and revenue.

Objective:

Build a robust ML-based delivery time prediction engine that:

- Accurately predicts delivery duration based on dynamic factors.
 - Learns from delivery history.
 - Adapts to city-wide operational behaviors.
 - Works in real-time via a deployable API.
-

3. Synthetic Dataset Creation: Simulating the Real World

Since real-world datasets from delivery platforms are protected by NDAs and privacy concerns, we developed a high-fidelity **synthetic dataset**. It mimics actual urban logistics operations using:

- Real city names (e.g., Dublin neighborhoods)
- Timestamps for order, dispatch, and delivery
- Calculated delivery distances (based on lat/lon pair logic)
- Vehicle types (Bike, Car, Van)
- Traffic levels (Low, Medium, High)
- Number of items

This dataset comprises **10,000 rows** with variability introduced through Gaussian noise, logical time lags, and event-driven behavior to make the data as realistic as possible.

4. Data Preprocessing & Feature Engineering

Key columns:

- `dispatch_delay_minutes`
- `delivery_distance_km`
- `order_hour`, `order_dayofweek`
- One-hot encoded features for `vehicle_type` and `traffic_level`

Outliers in `actual_duration_minutes` were removed using IQR filtering.

Feature Scaling:

All numeric variables were standardized using `StandardScaler` for optimal model convergence.

5. Modeling: From Baseline to Boosted Ensembles

We began with a **Linear Regression** model to set a baseline. Results:

- MAE: ~2.9 mins
- RMSE: ~3.8 mins
- R² Score: ~0.80

Then we moved to tree-based models:

- **Random Forest**
- **Gradient Boosting (XGBoost / GBRegressor)**

Performance (after tuning):

- RMSE: **3.63 mins**
- R² Score: **0.82**
- MAE: **2.80 mins**

These models were tuned using **GridSearchCV** for parameters like `n_estimators`, `learning_rate`, `max_depth`, and `subsample`.

6. Feature Importance Analysis

Top influential features:

- `delivery_distance_km`
- `dispatch_delay_minutes`
- `vehicle_type_Car`
- `traffic_level_High`
- `order_hour`

These insights help businesses **make data-driven dispatching decisions**.

7. Deployment: FastAPI + Ngrok in Colab

Once the model was finalized, it was serialized with `joblib`. The API was built using **FastAPI**, a high-performance Python web framework.

Endpoint:

`POST /predict`

Input: JSON with delivery context Output: Predicted delivery duration (in minutes)

Hosting:

To simulate production, the app was hosted using **ngrok** and integrated into Google Colab for portability.

8. Impact Evaluation

We evaluated business impact based on these KPIs:

- **Delivery Efficiency:** Better vehicle-time matching
 - **Route Planning:** Accurate ETA = smarter rerouting
 - **SLA Compliance:** Fewer penalties, higher CSAT
 - **Operational Costs:** Reduction in idle delivery agents
-

9. Real-World Use Case: Domino's, Zomato, Swiggy

Imagine Domino's Pizza using this:

- Before order acceptance, the model predicts delivery time.
- Dispatch team sees if bike vs car vs van changes ETA.
- Customers get dynamic ETAs (not static 30-minute defaults).

Swiggy or Zomato can use it to:

- Allocate drivers based on predicted efficiency
 - Auto-reschedule late orders
 - Feed data into **live ETA maps** for customers
-

10. Why This Project Stands Out

Unlike typical ML projects, this one is:

- **End-to-end deployed** (not just a notebook)
- **Backed by realistic synthetic data**
- **Industry-relevant**: Applies directly to last-mile delivery ops
- **Modular**: Easy to integrate with mobile/web front-ends
- **Explainable**: Feature importances make decisions transparent

Most importantly, it's **scalable**. The model can be trained on real data and moved to Azure, GCP, or AWS with ease.

11. Final Thoughts & Next Steps

This project proves that smart delivery time prediction is achievable without massive investment or private data.

Next enhancements:

- Add **route optimization layer** using OpenStreetMap or Google Directions API
 - Integrate **weather data** as a feature
 - Build a **Streamlit-based dashboard** for visual monitoring
 - Migrate to a **persistent cloud host** like Render or EC2
-

12. TL;DR

We built a predictive engine to estimate delivery duration using ML, deployed as a real-time API. It's modular, accurate, and made for real business adoption. Companies like Domino's and Zomato can benefit from it right away.

Want access to the code or API? Reach out or fork the [GitHub Repo](#) (only sample data included).

[Contact Me | LinkedIn | GitHub | Medium]

© 2025. Built with  by [Your Name].